

Image Classifier Project

August 1, 2019

1 Developing an AI application

Going forward, AI algorithms will be incorporated into more and more everyday applications. For example, you might want to include an image classifier in a smart phone app. To do this, you'd use a deep learning model trained on hundreds of thousands of images as part of the overall application architecture. A large part of software development in the future will be using these types of models as common parts of applications.

In this project, you'll train an image classifier to recognize different species of flowers. You can imagine using something like this in a phone app that tells you the name of the flower your camera is looking at. In practice you'd train this classifier, then export it for use in your application. We'll be using [this dataset](#) of 102 flower categories, you can see a few examples below.

The project is broken down into multiple steps:

- Load and preprocess the image dataset
- Train the image classifier on your dataset
- Use the trained classifier to predict image content

We'll lead you through each part which you'll implement in Python.

When you've completed this project, you'll have an application that can be trained on any set of labeled images. Here your network will be learning about flowers and end up as a command line application. But, what you do with your new skills depends on your imagination and effort in building a dataset. For example, imagine an app where you take a picture of a car, it tells you what the make and model is, then looks up information about it. Go build your own dataset and make something new.

First up is importing the packages you'll need. It's good practice to keep all the imports at the beginning of your code. As you work through this notebook and find you need to import a package, make sure to add the import up here.

```
In [1]: # Imports here
import torch
from torch import nn
from torch import optim
import torch.nn.functional as F
from torchvision import datasets, transforms, models
```

1.1 Load the data

Here you'll use `torchvision` to load the data ([documentation](#)). The data should be included alongside this notebook, otherwise you can [download it here](#). The dataset is split into three parts, training, validation, and testing. For the training, you'll want to apply transformations such as random scaling, cropping, and flipping. This will help the network generalize leading to better performance. You'll also need to make sure the input data is resized to 224x224 pixels as required by the pre-trained networks.

The validation and testing sets are used to measure the model's performance on data it hasn't seen yet. For this you don't want any scaling or rotation transformations, but you'll need to resize then crop the images to the appropriate size.

The pre-trained networks you'll use were trained on the ImageNet dataset where each color channel was normalized separately. For all three sets you'll need to normalize the means and standard deviations of the images to what the network expects. For the means, it's [0.485, 0.456, 0.406] and for the standard deviations [0.229, 0.224, 0.225], calculated from the ImageNet images. These values will shift each color channel to be centered at 0 and range from -1 to 1.

```
In [2]: data_dir = 'flowers'
        train_dir = data_dir + '/train'
        valid_dir = data_dir + '/valid'
        #test_dir = data_dir + '/test'
        test_dir = data_dir + '/test'

In [3]: # TODO: Define your transforms for the training, validation, and testing sets
        data_transforms = transforms.Compose([transforms.RandomRotation(30), transforms.RandomSize
                                              (), transforms.ToTensor(), transforms.Normalize([0.485
                                              0.456, 0.406], [0.229, 0.224, 0.225])])
        test_transform = transforms.Compose([transforms.Resize(225), transforms.CenterCrop(224), t
                                              transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0

        # TODO: Load the datasets with ImageFolder
        train_datasets = datasets.ImageFolder(train_dir, transform=data_transforms)
        test_data = datasets.ImageFolder(test_dir, transform=test_transform)
        validation_data = datasets.ImageFolder(valid_dir, transform=test_transform)

        # TODO: Using the image datasets and the trainforms, define the dataloaders
        trainLoader = torch.utils.data.DataLoader(train_datasets, batch_size=64, shuffle=True)
        testLoader = torch.utils.data.DataLoader(test_data, batch_size=64, shuffle=True)
        validationLoader = torch.utils.data.DataLoader(validation_data, batch_size=64, shuffle=True)
```

/opt/conda/lib/python3.6/site-packages/torchvision-0.2.1-py3.6.egg/torchvision/transforms/transf

1.1.1 Label mapping

You'll also need to load in a mapping from category label to category name. You can find this in the file `cat_to_name.json`. It's a JSON object which you can read in with the [json module](#). This will give you a dictionary mapping the integer encoded categories to the actual names of the flowers.

```
In [4]: import json

        with open('cat_to_name.json', 'r') as f:
            cat_to_name = json.load(f)
            #print(cat_to_name)
        with open('cat_to_name.json') as f:
            for val in f:
                print(val)

{"21": "fire lily", "3": "canterbury bells", "45": "bolero deep blue", "1": "pink primrose", "34": "lilac blue", "2": "orchid", "5": "red poppy", "7": "dandelion", "4": "lily of the valley", "6": "carnation", "10": "sunflower", "12": "iris", "32": "petunia", "14": "hyacinth", "16": "bluebell", "18": "tulip", "20": "daisy", "22": "marigold", "24": "geranium", "26": "fuchsia", "28": "hibiscus", "30": "lavender", "36": "roses", "40": "columbine", "42": "snapdragon", "44": "pansy", "46": "gerbil", "48": "hamster", "50": "squirrel", "52": "chipmunk", "54": "marmoset", "56": "lion", "58": "tiger", "60": "lioness", "62": "tigeress", "64": "lioness", "66": "tigeress", "68": "lioness", "70": "tigeress", "72": "lioness", "74": "tigeress", "76": "lioness", "78": "tigeress", "80": "lioness", "82": "tigeress", "84": "lioness", "86": "tigeress", "88": "lioness", "90": "tigeress", "92": "lioness", "94": "tigeress", "96": "lioness", "98": "tigeress", "100": "lioness"}
```

2 Building and training the classifier

Now that the data is ready, it's time to build and train the classifier. As usual, you should use one of the pretrained models from `torchvision.models` to get the image features. Build and train a new feed-forward classifier using those features.

We're going to leave this part up to you. Refer to [the rubric](#) for guidance on successfully completing this section. Things you'll need to do:

- Load a [pre-trained network](#) (If you need a starting point, the VGG networks work great and are straightforward to use)
- Define a new, untrained feed-forward network as a classifier, using ReLU activations and dropout
- Train the classifier layers using backpropagation using the pre-trained network to get the features
- Track the loss and accuracy on the validation set to determine the best hyperparameters

We've left a cell open for you below, but use as many as you need. Our advice is to break the problem up into smaller parts you can run separately. Check that each part is doing what you expect, then move on to the next. You'll likely find that as you work through each part, you'll need to go back and modify your previous code. This is totally normal!

When training make sure you're updating only the weights of the feed-forward network. You should be able to get the validation accuracy above 70% if you build everything right. Make sure to try different hyperparameters (learning rate, units in the classifier, epochs, etc) to find the best model. Save those hyperparameters to use as default values in the next part of the project.

One last important tip if you're using the workspace to run your code: To avoid having your workspace disconnect during the long-running tasks in this notebook, please read in the earlier page in this lesson called Intro to GPU Workspaces about Keeping Your Session Active. You'll want to include code from the `workspace_utils.py` module.

Note for Workspace users: If your network is over 1 GB when saved as a checkpoint, there might be issues with saving backups in your workspace. Typically this happens with wide dense layers after the convolutional layers. If your saved checkpoint is larger than 1 GB (you can open a terminal and check with `ls -lh`), you should reduce the size of your hidden layers and train again.

```
In [5]: from torchvision import models, datasets
```

```
In [6]: #choose = ['models.resnet50(pretrained=True)', 'models.GoogLeNet(pretrained=True)', 'models.vgg19(pretrained=True)']
        #choose1 = models.resnet50(pretrained=True)
```

```
In [7]: model = models.vgg19(pretrained=True)
```

```
        #print(model)
```

```
Downloading: "https://download.pytorch.org/models/vgg19-dcbb9e9d.pth" to /root/.torch/models/vgg19-dcbb9e9d.pth
100%|| 574673361/574673361 [00:09<00:00, 61850946.81it/s]
```

```
In [8]: #model = models.vgg16(pretrained=True)
        #print(model)
```

```
In [9]: # Freeze parameters so we don't backprop through them
        for param in model.parameters():
            param.requires_grad = False
```

```
In [10]: model.classifier = nn.Sequential(nn.Linear(25088,4096),nn.ReLU(),nn.Dropout(0.3),nn.Linear(4096,1000),nn.Dropout(0.3),nn.Linear(250,102),nn.LogSoftmax(dim=1))
```

```
In [11]: device = torch.device('cpu')
        model.to(device)
```

```
Out[11]: VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace)
    (16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (17): ReLU(inplace)
    (18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```

(22): ReLU(inplace)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU()
  (2): Dropout(p=0.3)
  (3): Linear(in_features=4096, out_features=1000, bias=True)
  (4): ReLU()
  (5): Dropout(p=0.3)
  (6): Linear(in_features=1000, out_features=250, bias=True)
  (7): ReLU()
  (8): Dropout(p=0.3)
  (9): Linear(in_features=250, out_features=102, bias=True)
  (10): LogSoftmax()
)
)

```

```

In [12]: criterion = nn.NLLLoss()
         optimizer = optim.Adam(model.classifier.parameters(),lr=0.001)

```

```

In [13]: from workspace_utils import active_session

```

```

In [ ]: Save the checkpoint

```

```

In [67]: classidx = train_datasets.class_to_idx
         def SaveCheckpoint(model,classidx,epchos,ModelName,path):
             model.class_to_idx = classidx
             model.cpu
             checkpoint = {'architecture': ModelName, 'classifier': model.classifier, 'class_to_idx':
                           'state_dict': model.state_dict(), 'epoch': epoch}

             torch.save(checkpoint, path)

```

```

In [ ]: print(device)
         Train_loss_graph = []

```

```

Test_loss_graph = []
epochs = 30
accuracytemp = 0
steps = 0
running_loss = 0
print_every = 35
for epoch in range(epochs):
    for inputs, labels in trainLoader:
        steps += 1
        # Move input and label tensors to the default device
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        logps = model.forward(inputs)
        loss = criterion(logps, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    if steps % print_every == 0:
        test_loss = 0
        accuracy = 0
        model.eval()
        with torch.no_grad():
            for inputs, labels in validationLoader:
                inputs, labels = inputs.to(device), labels.to(device)
                logps = model.forward(inputs)
                batch_loss = criterion(logps, labels)

                test_loss += batch_loss.item()

                # Calculate accuracy
                ps = torch.exp(logps)####bueatify result
                top_p, top_class = ps.topk(1, dim=1)
                equals = top_class == labels.view(*top_class.shape)
                accuracy += torch.mean(equals.type(torch.FloatTensor)).item()

            #####

        print(f"Epoch {epoch+1}/{epochs}.. "
              f"Train loss: {running_loss/print_every:.3f}.. "
              f"Test loss: {test_loss/len(validationLoader):.3f}.. "
              f"Test accuracy: {accuracy/len(validationLoader):.3f}")
        accuracyCal = (accuracy/len(validationLoader))

```

```

        if accuracyCal > accuracytemp:
            accuracytemp=accuracyCal
            SaveCheckpoint(model,classidx,epoch,ModelName='vgg19',path='my_m
            print("*****Save executed *****")

    Train_loss_graph.append(running_loss)
    Test_loss_graph.append(test_loss)
    running_loss = 0
    model.train()

cuda
Epoch 1/30.. Train loss: 1.267.. Test loss: 0.735.. Test accuracy: 0.817
*****Save executed *****
Epoch 1/30.. Train loss: 1.242.. Test loss: 0.727.. Test accuracy: 0.805
Epoch 2/30.. Train loss: 1.308.. Test loss: 0.756.. Test accuracy: 0.788
Epoch 2/30.. Train loss: 1.205.. Test loss: 0.704.. Test accuracy: 0.800
Epoch 2/30.. Train loss: 1.199.. Test loss: 0.682.. Test accuracy: 0.816
Epoch 3/30.. Train loss: 1.189.. Test loss: 0.642.. Test accuracy: 0.824
*****Save executed *****
Epoch 3/30.. Train loss: 1.256.. Test loss: 0.628.. Test accuracy: 0.818
Epoch 3/30.. Train loss: 1.179.. Test loss: 0.763.. Test accuracy: 0.803
Epoch 4/30.. Train loss: 1.156.. Test loss: 0.656.. Test accuracy: 0.811
Epoch 4/30.. Train loss: 1.079.. Test loss: 0.644.. Test accuracy: 0.825
*****Save executed *****
Epoch 4/30.. Train loss: 1.097.. Test loss: 0.698.. Test accuracy: 0.814
Epoch 5/30.. Train loss: 1.188.. Test loss: 0.646.. Test accuracy: 0.824
Epoch 5/30.. Train loss: 1.030.. Test loss: 0.580.. Test accuracy: 0.835
*****Save executed *****
Epoch 5/30.. Train loss: 1.079.. Test loss: 0.605.. Test accuracy: 0.848
*****Save executed *****
Epoch 6/30.. Train loss: 1.128.. Test loss: 0.615.. Test accuracy: 0.843
Epoch 6/30.. Train loss: 0.997.. Test loss: 0.553.. Test accuracy: 0.855
*****Save executed *****
Epoch 6/30.. Train loss: 1.069.. Test loss: 0.570.. Test accuracy: 0.849
Epoch 7/30.. Train loss: 1.054.. Test loss: 0.610.. Test accuracy: 0.844
Epoch 7/30.. Train loss: 1.041.. Test loss: 0.608.. Test accuracy: 0.847
Epoch 7/30.. Train loss: 1.042.. Test loss: 0.585.. Test accuracy: 0.841
Epoch 8/30.. Train loss: 1.029.. Test loss: 0.641.. Test accuracy: 0.827
Epoch 8/30.. Train loss: 1.018.. Test loss: 0.637.. Test accuracy: 0.836
Epoch 8/30.. Train loss: 0.985.. Test loss: 0.619.. Test accuracy: 0.830
Epoch 9/30.. Train loss: 1.014.. Test loss: 0.502.. Test accuracy: 0.865
*****Save executed *****
Epoch 9/30.. Train loss: 1.002.. Test loss: 0.515.. Test accuracy: 0.865
Epoch 9/30.. Train loss: 0.945.. Test loss: 0.521.. Test accuracy: 0.862
Epoch 10/30.. Train loss: 0.942.. Test loss: 0.560.. Test accuracy: 0.858
Epoch 10/30.. Train loss: 0.906.. Test loss: 0.557.. Test accuracy: 0.855
Epoch 10/30.. Train loss: 0.972.. Test loss: 0.540.. Test accuracy: 0.851

```

```

Epoch 11/30.. Train loss: 0.932.. Test loss: 0.506.. Test accuracy: 0.867
*****Save executed *****
Epoch 11/30.. Train loss: 0.943.. Test loss: 0.538.. Test accuracy: 0.864
Epoch 11/30.. Train loss: 0.996.. Test loss: 0.524.. Test accuracy: 0.866
Epoch 12/30.. Train loss: 0.970.. Test loss: 0.541.. Test accuracy: 0.860
Epoch 12/30.. Train loss: 0.900.. Test loss: 0.528.. Test accuracy: 0.865
Epoch 12/30.. Train loss: 0.865.. Test loss: 0.482.. Test accuracy: 0.866
Epoch 13/30.. Train loss: 0.965.. Test loss: 0.546.. Test accuracy: 0.863
Epoch 13/30.. Train loss: 0.923.. Test loss: 0.561.. Test accuracy: 0.840
Epoch 13/30.. Train loss: 0.887.. Test loss: 0.591.. Test accuracy: 0.844
Epoch 14/30.. Train loss: 0.980.. Test loss: 0.561.. Test accuracy: 0.855
Epoch 14/30.. Train loss: 0.883.. Test loss: 0.542.. Test accuracy: 0.860
Epoch 14/30.. Train loss: 0.863.. Test loss: 0.544.. Test accuracy: 0.848
Epoch 15/30.. Train loss: 0.982.. Test loss: 0.533.. Test accuracy: 0.865
Epoch 15/30.. Train loss: 0.914.. Test loss: 0.498.. Test accuracy: 0.871
*****Save executed *****
Epoch 15/30.. Train loss: 0.951.. Test loss: 0.602.. Test accuracy: 0.848
Epoch 16/30.. Train loss: 0.921.. Test loss: 0.489.. Test accuracy: 0.876
*****Save executed *****
Epoch 16/30.. Train loss: 0.927.. Test loss: 0.535.. Test accuracy: 0.849
Epoch 16/30.. Train loss: 0.900.. Test loss: 0.500.. Test accuracy: 0.870
Epoch 17/30.. Train loss: 0.830.. Test loss: 0.517.. Test accuracy: 0.865
Epoch 17/30.. Train loss: 0.919.. Test loss: 0.497.. Test accuracy: 0.872
Epoch 17/30.. Train loss: 0.883.. Test loss: 0.465.. Test accuracy: 0.875
Epoch 18/30.. Train loss: 0.838.. Test loss: 0.448.. Test accuracy: 0.889
*****Save executed *****
Epoch 18/30.. Train loss: 0.814.. Test loss: 0.524.. Test accuracy: 0.876
Epoch 19/30.. Train loss: 0.871.. Test loss: 0.496.. Test accuracy: 0.874
Epoch 19/30.. Train loss: 0.862.. Test loss: 0.525.. Test accuracy: 0.865
Epoch 19/30.. Train loss: 0.778.. Test loss: 0.483.. Test accuracy: 0.879
Epoch 20/30.. Train loss: 0.830.. Test loss: 0.522.. Test accuracy: 0.866
Epoch 20/30.. Train loss: 0.822.. Test loss: 0.520.. Test accuracy: 0.868
Epoch 20/30.. Train loss: 0.912.. Test loss: 0.508.. Test accuracy: 0.869
Epoch 21/30.. Train loss: 0.956.. Test loss: 0.551.. Test accuracy: 0.857
Epoch 21/30.. Train loss: 0.874.. Test loss: 0.500.. Test accuracy: 0.878

```

2.1 Testing your network

It's good practice to test your trained network on test data, images the network has never seen either in training or validation. This will give you a good estimate for the model's performance on completely new images. Run the test images through the network and measure the accuracy, the same way you did validation. You should be able to reach around 70% accuracy on the test set if the model has been trained well.

```
In [ ]:
```

```
In [12]: #Testing
         #model.eval()
```



```
In [14]: model.name = "vgg19"
```

```
In [15]: checkpoint = {'architecture':model.name, 'classifier':model.classifier, 'class_to_idx':mo
            'state_dict':model.state_dict(), 'epoch':15}
```

```
In [16]: torch.save(checkpoint, 'my_model.pth')
```

2.3 Loading the checkpoint

At this point it's good to write a function that can load a checkpoint and rebuild the model. That way you can come back to this project and keep working on it without having to retrain the network.

```
In [17]: # TODO: Write a function that loads a checkpoint and rebuilds the model
```

```
def load_checkpoint_cpu():
    checkpoint = torch.load('my_model.pth')
    model = models.vgg19(pretrained=True)
    for param in model.parameters(): param.requires_grad = False
    model.class_to_idx = checkpoint['class_to_idx']
    model.classifier = checkpoint['classifier']
    model.load_state_dict(checkpoint['state_dict'])
    return model
```

```
In [18]: model = load_checkpoint_cpu()
        criterion = nn.NLLLoss()
        optimizer = optim.Adam(model.classifier.parameters(), lr=0.001)
```

3 Inference for classification

Now you'll write a function to use a trained network for inference. That is, you'll pass an image into the network and predict the class of the flower in the image. Write a function called `predict` that takes an image and a model, then returns the top *K* most likely classes along with the probabilities. It should look like

```
probs, classes = predict(image_path, model)
print(probs)
print(classes)
> [ 0.01558163  0.01541934  0.01452626  0.01443549  0.01407339]
> ['70', '3', '45', '62', '55']
```

First you'll need to handle processing the input image such that it can be used in your network.

3.1 Image Preprocessing

You'll want to use PIL to load the image ([documentation](#)). It's best to write a function that preprocesses the image so it can be used as input for the model. This function should process the images in the same manner used for training.

First, resize the images where the shortest side is 256 pixels, keeping the aspect ratio. This can be done with the `thumbnail` or `resize` methods. Then you'll need to crop out the center 224x224 portion of the image.

Color channels of images are typically encoded as integers 0-255, but the model expects floats 0-1. You'll need to convert the values. It's easiest with a Numpy array, which you can get from a PIL image like so `np_image = np.array(pil_image)`.

As before, the network expects the images to be normalized in a specific way. For the means, it's [0.485, 0.456, 0.406] and for the standard deviations [0.229, 0.224, 0.225]. You'll want to subtract the means from each color channel, then divide by the standard deviation.

And finally, PyTorch expects the color channel to be the first dimension but it's the third dimension in the PIL image and Numpy array. You can reorder dimensions using `ndarray.transpose`. The color channel needs to be first and retain the order of the other two dimensions.

```
In [19]: import PIL
         from PIL import Image
         import numpy as np
         import matplotlib.pyplot as plt

         import matplotlib.pyplot as plt

In [20]: def process_image(image):
         ''' Scales, crops, and normalizes a PIL image for a PyTorch model,
             returns an Numpy array
             '''

         pilImg = Image.open(f'{image}' + '.jpg')
         TransformImgFeeder = transforms.Compose([transforms.Resize(226), transforms.CenterCrop(224),
                                                    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])
         pic_tensor = TransformImgFeeder(pilImg)
         numpyArray = np.array(pic_tensor)
         return numpyArray

         # TODO: Process a PIL image for use in a PyTorch model
```

To check your work, the function below converts a PyTorch tensor and displays it in the notebook. If your `process_image` function works, running the output through this function should return the original image (except for the cropped out portions).

```
In [21]: def imshow(image, ax=None, title=None):
         """Imshow for Tensor."""
         if ax is None:
             fig, ax = plt.subplots()

         # PyTorch tensors assume the color channel is the first dimension
             # but matplotlib assumes is the third dimension
             #image = image.numpy().transpose((1, 2, 0))

         image = image.transpose((1, 2, 0))
```

```

# Undo preprocessing
mean = np.array([0.485, 0.456, 0.406])
std = np.array([0.229, 0.224, 0.225])
image = std * image + mean

# Image needs to be clipped between 0 and 1 or it looks like noise when displayed
image = np.clip(image, 0, 1)

ax.imshow(image)

return ax

```

```

In [22]: image_graph = test_dir+'10/'+image_07090'
        img_test_o = process_image(image_graph)

```

```


In [23]: imshow(img_test_o, ax=None, title=None)

```

```

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9862ec8dd8>

```



3.2 Class Prediction

Once you can get images in the correct format, it's time to write a function for making predictions with your model. A common practice is to predict the top 5 or so (usually called top- K) most probable classes. You'll want to calculate the class probabilities then find the K largest values.

To get the top K largest values in a tensor use `x.topk(k)`. This method returns both the highest k probabilities and the indices of those probabilities corresponding to the classes. You need to

convert from these indices to the actual class labels using `class_to_idx` which hopefully you added to the model or from an `ImageFolder` you used to load the data (Section 2.2). Make sure to invert the dictionary so you get a mapping from index to class as well.

Again, this method should take a path to an image and a model checkpoint, then return the probabilities and classes.

```
probs, classes = predict(image_path, model)
print(probs)
print(classes)
> [ 0.01558163  0.01541934  0.01452626  0.01443549  0.01407339]
> ['70', '3', '45', '62', '55']

In [26]: def predict(image_path, model, top_k = 5):
        model.to('cpu')
        model.eval()
        image = process_image(image_path)
        #pytorch tensor = torch.tensor(image)

        image_tensor = torch.from_numpy(image).type(torch.FloatTensor)

        pytorch_tensor = image_tensor.unsqueeze(0) #add a 1 as the first argument of our tensor
        output = model.forward(pytorch_tensor)
        output1 = torch.exp(output) #the predicted probability
        prob, indices = output1.topk(top_k)
        top_probs = prob.detach().numpy().tolist()[0]
        top_labs = indices.detach().numpy().tolist()[0]
        #[dict[model.class_to_idx] for model.class_to_idx in indices ]
        #map(model.class_to_idx.get, indices)
        lis=[]
        #lis = Null
        idx_to_class = {val: key for key, val in model.class_to_idx.items()}
        #print(idx_to_class)
        for i in top_labs:
            #val = str(i)
            lis.append(idx_to_class[i])
        #print("list values = ", lis)
        return top_probs, lis
```

```
In [27]: prob , classes = predict(image_graph,model)
```

```
In [28]: print(prob)
         print(classes)
```

```
[1.0, 3.872737200372178e-12, 1.0309890555214654e-13, 1.8254737389126856e-14, 3.204654076488316e-15]
['10', '22', '92', '14', '38']
```

3.3 Sanity Checking

Now that you can use a trained model for predictions, check to make sure it makes sense. Even if the testing accuracy is high, it's always good to check that there aren't obvious bugs. Use matplotlib to plot the probabilities for the top 5 classes as a bar graph, along with the input image. It should look like this:

You can convert from the class integer encoding to actual flower names with the `cat_to_name.json` file (should have been loaded earlier in the notebook). To show a PyTorch tensor as an image, use the `imshow` function defined above.

```
In [29]: with open('cat_to_name.json', 'r') as f:
          cat_to_name = json.load(f)
          print(cat_to_name)
```

```
{'21': 'fire lily', '3': 'canterbury bells', '45': 'bolero deep blue', '1': 'pink primrose', '34': 'black-cherry lily', '27': 'corn poppy', '51': 'carnation', '18': 'catenula', '9': 'sea purslane', '30': 'sunflower', '65': 'iris-violet blue', '24': 'gladiolus', '42': 'iris-violet white', '26': 'iris-pink', '39': 'iris-white', '7': 'iris-black', '64': 'iris-blue', '2': 'iris-yellow', '54': 'iris-white', '8': 'iris-black', '63': 'iris-blue', '25': 'iris-yellow', '53': 'iris-white', '62': 'iris-blue', '23': 'iris-yellow', '52': 'iris-white', '61': 'iris-blue', '22': 'iris-yellow', '50': 'iris-white', '60': 'iris-blue', '20': 'iris-yellow', '49': 'iris-white', '59': 'iris-blue', '19': 'iris-yellow', '48': 'iris-white', '58': 'iris-blue', '17': 'iris-yellow', '47': 'iris-white', '57': 'iris-blue', '16': 'iris-yellow', '46': 'iris-white', '56': 'iris-blue', '15': 'iris-yellow', '45': 'iris-white', '55': 'iris-blue', '14': 'iris-yellow', '44': 'iris-white', '54': 'iris-blue', '13': 'iris-yellow', '43': 'iris-white', '53': 'iris-blue', '12': 'iris-yellow', '42': 'iris-white', '52': 'iris-blue', '11': 'iris-yellow', '41': 'iris-white', '51': 'iris-blue', '10': 'iris-yellow', '40': 'iris-white', '50': 'iris-blue', '9': 'iris-yellow', '39': 'iris-white', '49': 'iris-blue', '8': 'iris-yellow', '38': 'iris-white', '48': 'iris-blue', '7': 'iris-yellow', '37': 'iris-white', '47': 'iris-blue', '6': 'iris-yellow', '36': 'iris-white', '46': 'iris-blue', '5': 'iris-yellow', '35': 'iris-white', '45': 'iris-blue', '4': 'iris-yellow', '34': 'iris-white', '44': 'iris-blue', '3': 'iris-yellow', '33': 'iris-white', '43': 'iris-blue', '2': 'iris-yellow', '32': 'iris-white', '42': 'iris-blue', '1': 'iris-yellow', '31': 'iris-white', '41': 'iris-blue', '0': 'iris-yellow', '30': 'iris-white', '40': 'iris-blue'}
```

```
In [ ]:
```

```
In [48]: def PredictProb(imagePath,model):
          ##lets get the image
          imageConv = process_image(imagePath)
          #imshow(imageConv, ax=None, title=None)
          #print("print val *****\n")
          topP, topC = predict(imagePath, model,top_k = 5)

          lis1 =[]
          for i in topC:
              val = str(i)
              lis1.append(cat_to_name[val])

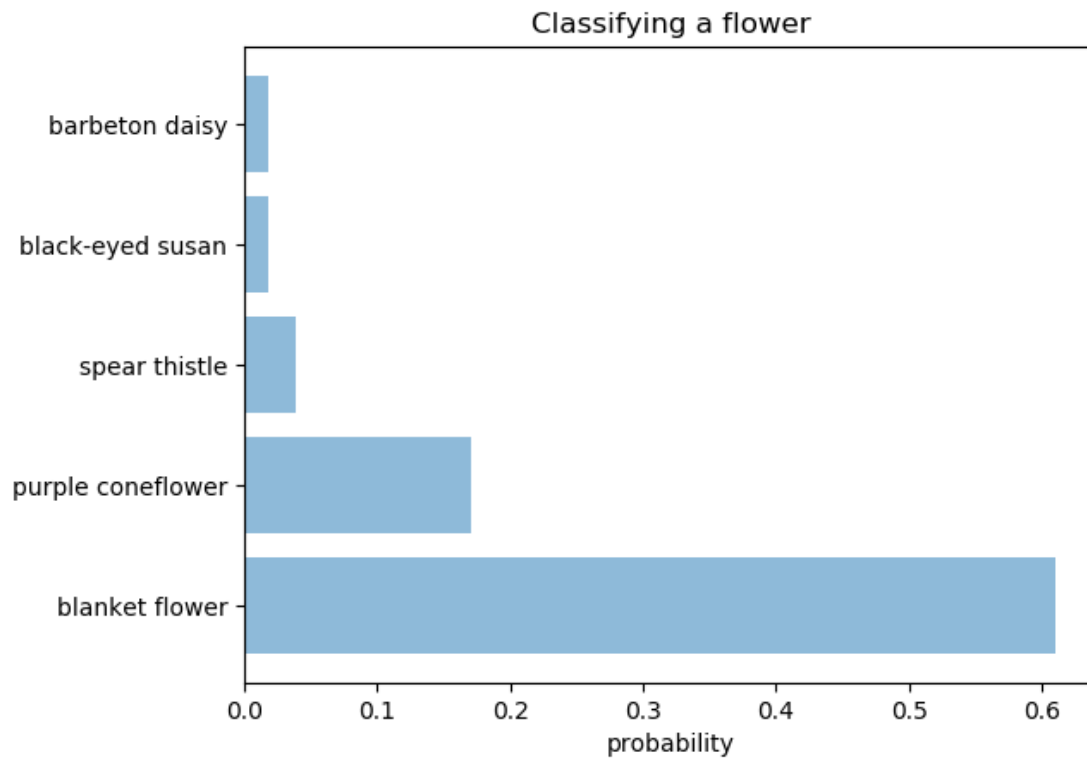
          objects = tuple(lis1)


          y_pos = np.arange(len(objects))

          plt.barh(y_pos, topP, align='center', alpha=0.5)
          plt.yticks(y_pos, objects)
          plt.xlabel('probability')
          plt.title('Classifying a flower')

          plt.show()
          def PredictProbPic(imagePath):
              ##lets get the image
              imageConv = process_image(imagePath)
              imshow(imageConv, ax=None, title=None)
              #imshow(imageConv, ax=None, title=None)
```

```
In [51]: imagePath = test_dir+'/100/'+image_07902'  
        PredictProb(imagePath,model)  
        PredictProbPic(imagePath)
```





```
In [ ]:
```

```
In [53]: !!jupyter nbconvert *.ipynb
```

```
Out[53]: ['[NbConvertApp] Converting notebook Image Classifier Project.ipynb to html',  
          '[NbConvertApp] Writing 737247 bytes to Image Classifier Project.html',  
          '[NbConvertApp] Converting notebook Image Classifier Project-zh.ipynb to html',  
          '[NbConvertApp] Writing 295106 bytes to Image Classifier Project-zh.html',  
          '[NbConvertApp] Converting notebook Temp.ipynb to html',  
          '[NbConvertApp] Writing 463687 bytes to Temp.html']
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```